



Assignment 2: Spark Data Analytics

Individual Work: 20%

10.05.2021

1 Introduction

This assignment tests your ability to use various Spark APIs to implement given workloads. It also test your ability to analyze and tune the performance of Spark application.

The assignment data set contains tweet objects downloaded from **Twitter** using **Tweeter standard search API**. The downloaded tweet objects are stored in a single JSON file. A tweet object may refer to a *general tweet*, a *retweet* or a *reply* to a tweet. A *general tweet* is “a message posted to Twitter containing text, photos, a GIF, and/or video”¹. A reply is a response to another person’s tweet². “A Retweet is a re-posting of a Tweet.”³ Below is an example of a retweet object:

```
1  {
2      "id": 1390088382659895296,
3      "created_at": "2021-05-05 23:37:51",
4      "text": "RT @nytimes: Breaking News: The Biden administration will
           support lifting patent protections for Covid-19 vaccines, a
           breakthrough for glob...",
5      "user_id": 17799542,
6      "retweet_id": 1390027514332991489,
7      "retweet_user_id": 807095,
8      "user_mentions": [
9          {
10             "id": 807095,
11             "indices": [
12                 3,
13                 11
14             ]
15         }
16     ]
17 }
```

¹About different types of Tweets General Tweets

²About different types of Tweets General Tweets

³Retweet FAQs

The common fields in all tweet objects are:

- `id`: the unique id of the tweet.
- `created_at`: the date and time the tweet is created.
- `text`: the textual content of the tweet
- `user_id`: the id of the user who created the tweet.
- `retweet_count`: retweet count of the tweet or its parent tweet if the tweet is a *retweet*
- `favorite_count`: favorite count of the tweet

The optional fields are:

- `retweet_id`: included only in *retweet* object. It contains the id of the tweet it re-posts.
- `retweet_user_id`: included only in *retweet* object. It contains the `user_id` of the tweet it re-posts.
- `replyto_id`: included only in *reply* object. It contains the id of the tweet it responds to.
- `replyto_userid`: included only in *reply* object. It contains the `user_id` of the tweet it responds to
- `user_mentions`: included only if the tweet text mentions one or more users. It is an array of `{id, indices}` showing the id of the users mentioned in the tweet text, as well as the location this user is mentioned.
- `hash_tags`: included only if the tweet text contains one or more hash_tags. It is an array of `{tag, indices}` showing the hash tag appearing in the tweet text and the location it appears

The data set is provided as a multi line JSON file and should be loaded into Spark data frame using the following statement:

```
spark.read.option("multiline", "true").json(file_path)
```

Missing fields in an object will have `null` value in the data frame.

2 Workload One

In this workload, you are asked to find the top 5 users with similar interest as a given user id. We measure a user's interest based on the tweets they have replied or retweeted. We can represent each user as a tokenized document, with all tokens as ids of the tweets the user has replied or retweeted. Assume we have a sample data set with five tweets: t_1, t_2, t_3, t_4, t_5 and five users: u_1, u_2, u_3, u_4, u_5 , table 1 shows the reply/retweet relation from user to tweet. The last column shows the corresponding document representation of each user. For instance, u_1 has retweeted t_3, t_5 and replied t_4, t_5 , the document representation is (t_3, t_4, t_5, t_5) .

Table 1: Sample Data

UID	TID					Document Representation
	t1	t2	t3	t4	t5	
u1			rt	rp	rt,rp	(t_3, t_4, t_5, t_5)
u2		rp	rt,rp	rp		(t_2, t_3, t_3, t_4)
u3	rt				rt,rp	(t_1, t_5, t_5)
u4	rt				rp	(t_1, t_5)
u5		rt			rp	(t_2, t_5)

To implement the workload, you need to transform each user document representation into a vector and use Cosine similarity to compute the similarity between two user vectors. SparkML provides a number of feature extractors such as TF-IDF, Word2Vec and more. You are asked to experiment with two feature extractors and compare the execution performance in local and EMR cluster (minimum 3 nodes) setting. You can use any user id in the data set as the given user id. For each feature extractor, you can set your own parameter values. There is no standard result to the workload.

3 Workload Two: User Recommendation

We consider two groups of users. One group contains users that appear in the `user_id` field of a tweet object, we call them *tweet users*; another group contains users mentioned in a tweet, we call them *mention users*. In this workload, you are asked to use collaborative filtering algorithm in SparkML library to recommend top 5 *mention users* to each *tweet user* in the data set.

Collaborative filtering aims to fill in the missing entries of a user-item association matrix. We treat *tweet users* as 'user' and *mention users* as 'item' in this setting. The rating is computed as number of times a *tweet user* mentions a *mention user*. You are asked to prepare the raw data in the format as required by the collaborative filter algorithm and build a model to perform the recommendation. Similar to workload one, you are asked to run the workload in local and EMR cluster setting and observe the performance statistics.

4 Coding and Execution Requirement

Below are requirements on coding and Execution:

- You should implement both workloads in **PySpark**. You can use SparkRDD, SparkSQL and SparkML API.
- There should be a single driver program for both workloads. You should provide python script version of the implementation as well as Spark submit script. There should be enough comment to help understand the data flow of each workload. You may also submit a notebook version.
- You should not modify the input data file in any way and your code should read the data file from the same directory as the source code.

5 Deliverable

There are two deliverables: **source code** and **project report**. Both are due on **Wednesday 26th of May 23:59 (Week 12)**. There will be separate submission links for source code and project report on Canvas. The report needs to be submitted as a pdf file. ~~The source code could be submitted as a zip file containing python script(s) and spark submission script(s) and/or Jupyter notebook file.~~ The source code could be submitted as a zip file containing Python script(s), Spark submission script(s) and, optionally, Jupyter Notebook files. In other words, the submission of Python script(s) and Spark submission script(s) are compulsory, while Jupyter notebook file is optional. Please do not include data file in the zip file. There will be penalty point for including data file.

There will be a 10 minutes demo in week 12, 13, or STUVAC. You need to attend the demo to receive mark for this assignment!

6 Report Structure

The report should contain four sections:

- A brief introduction
- Workload one
 - design
 - performance analysis
- Workload two
 - design

- performance analysis
- A brief conclusion

The design section should describe overall data flow of each workload. You can include DAG chart to visualize the sequence of operations. You may include short code snippets in this section to help explain certain points. You should highlight performance optimization features you use in the implementation; this could include choosing certain operations to minimize shuffling, using RDD persistence or other features.

The Performance Analysis section should give an overview of the execution environment as well as the execution configuration (executor capacity, number of executors, etc) you have chosen. It should include a few graphs and/or tables highlighting performance variations under various settings. You can include screenshots of history server pages. You should also include enough explanation of the performance statistics and differences observed.

7 Mark Distribution

- Implementation/Source code: 14 points
 - 7 points for each workload
 - Your code should be executable and follow the coding requirement to receive points for **coding** section.
- Report: 6 points
 - The report (including all headings, labels and legends in figures and graphs, and captions) must be written in Academic English. All sources that you have used must be properly cited and referenced.
 - Your report must be submitted in the submission links for report. Otherwise, you may lose mark for that.
 - Your code should be executable and follow the coding requirement to receive points for **report** section.